

Parallelization Report

Max Stritzinger

For this project, I needed to discretize the heat equation to steady state. This program considers a heat source function, $S(x,y)$, boundary conditions, and calculates the heat at given points.

I first implemented this function serially. The code is given as twoD.cpp. The source function used is $S(x,y) = -\sin(x) - \sin(y)$, however any function may be used. This makes the exact solution, $u(x,y) = \sin(x) + \sin(y)$. I use the given equation and the boundary conditions to then iteratively calculate predicted values of the exact solution until we reach a specified error threshold between iterations or a maximum number of iterations. In our solution, I used a tolerance of $1E-15$.

I then tested my code comparing the output predicted values with the exact solution to ensure that our approximation was correct.

I then attempted to parallelize the code. There are several loops in my serial implementation that are targets for parallelization. The most obvious is where all values of U^{n+1} are approximated. This was a successful parallelization that reduced runtimes for all m,n that were greater than 50.

There are two other candidate loops for parallelization. One is the calculation of the error between U^n and U^{n+1} . I attempted parallelizing first using a parallel for with a critical block for the update. This critical section slowed the code down so much it could not finish 100×100 on two threads. I then tried using an atomic update, however this too slowed down the code compared to the serial version. Our last attempt was creating a $m \times n$ error grid where each error value was calculated in parallel and then we looked for the max error serially. This solution only slowed down the execution by $\sim .2$ seconds for a 100×100 grid. It's predicted that this solution would actually be faster for very large values of m,n . (given in twoD_parallel2.cpp).



The other loop that could be parallelized is the update loop, however a quick test showed it was fastest to just do this serially.

I then ran timed test of the parallel code to see how numbers of threads affected runtime, speedup, and efficiency. These tests were run for 100×100 and 150×150 grids using the icpc compiler on a 4 core processor. The results table is included as results.xlsx.

Speedup was calculated as $-\frac{(Time_{n-1} - Time_n)}{(NumThreads_{n-1} - NumThreads_n)}$. For both 100×100 and 150×150 the highest speedup came between going from one thread to 2 threads. Speedup increased up until 4 threads after which additional threads slowed the program down.

Efficiency was calculated as $\log \frac{1}{Time * NumThreads} + c$ where c is a constant used to normalize the values to be positive. The maximum efficiency came from using a single thread and decreased with each additional thread.

Graphs shown below.



Efficiency vs Num Threads

