

## Parallelization Report

Max Stritzinger

For this project, I needed to discretize the heat equation to steady state. This program considers a heat source function,  $S(x,y)$ , boundary conditions, and calculates the heat at given points.

I first implemented this function serially. The code is given as twoD.cpp. The source function used is  $S(x,y) = -\sin(x) - \sin(y)$ , however any function may be used. This makes the exact solution,  $u(x,y) = \sin(x) + \sin(y)$ . I use the given equation and the boundary conditions to then iteratively calculate predicted values of the exact solution until we reach a specified error threshold between iterations or a maximum number of iterations. In our solution, I used a tolerance of  $1E-15$ .

I then tested my code comparing the output predicted values with the exact solution to ensure that our approximation was correct.

I then attempted to parallelize the code. There are several loops in my serial implementation that are targets for parallelization. The most obvious is where all values of  $U^{n+1}$  are approximated. This was a successful parallelization that reduced runtimes for all  $m,n$  that were greater than 50.

There are two other candidate loops for parallelization. One is the calculation of the error between  $U^n$  and  $U^{n+1}$ . I attempted parallelizing first using a parallel for with a critical block for the update. This critical section slowed the code down so much it could not finish  $100 \times 100$  on two threads. I then tried using an atomic update, however this too slowed down the code compared to the serial version. Our last attempt was creating a  $m \times n$  error grid where each error value was calculated in parallel and then we looked for the max error serially. This solution only slowed down the execution by  $\sim .2$  seconds for a  $100 \times 100$  grid. It's predicted that this solution would actually be faster for very large values of  $m,n$ . (given in twoD\_parallel2.cpp).

The other loop that could be parallelized is the update loop, however a quick test showed it was fastest to just do this serially.

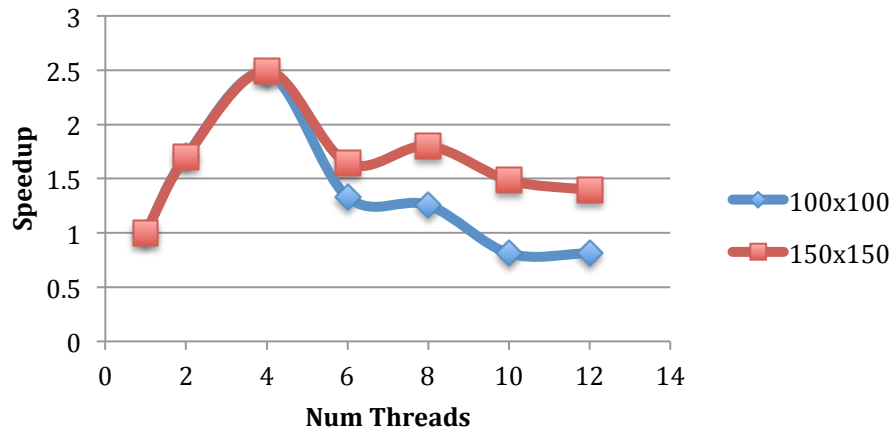
I then ran timed test of the parallel code to see how numbers of threads affected runtime, speedup, and efficiency. These tests were run for  $100 \times 100$  and  $150 \times 150$  grids using the icpc compiler on a 4 core processor. The results table is included as results.xlsx.

For both  $100 \times 100$  and  $150 \times 150$  the highest speedup came using the number of threads = number of cores = 4. For  $100 \times 100$ , speedup was  $> 1$  until 8 cores, while for  $150 \times 150$  it was  $> 1$  up until 20.


The maximum efficiency came from using a single thread and decreased with each additional thread.

Graphs shown below.

## Speedup vs Num Threads



## Time vs Num Threads



## Efficiency vs Num Threads

